

## **EXPERIMENT-03**

- **AIM:** - To design a normalized relational database schema using SQL to manage students, courses, and their enrollments with proper integrity constraints; perform data insertion, transactions using SAVEPOINTS and ROLLBACK, simulate error handling with faulty records, and retrieve meaningful reports using joins to display students' enrollment and grade details.
- **THEORY:** -
  1. **Relational Database Design & Normalization:** -A relational database stores data in structured tables (relations) where each table represents an entity (e.g., Students, Courses). To avoid redundancy and maintain data integrity, databases are normalized — i.e., broken down into multiple related tables that satisfy normal forms (up to 3NF in this case). The key entities in this schema are:

Students (student\_id, name, dob)

Courses (course\_id, title)

Enrollments (enroll\_id, student\_id, course\_id, grade)

Here, Enrollments act as a junction table to manage many-to-many relationships between Students and Courses.

### **2. Primary and Foreign Key Constraints**

Primary Keys (PK) uniquely identify each record in a table (e.g., student\_id, course\_id, enroll\_id).

Foreign Keys (FK) maintain referential integrity, ensuring that data in one table corresponds to valid data in another (e.g., student\_id in Enrollments must exist in Students).

### **3. SQL Transactions**

A transaction is a sequence of one or more SQL operations treated as a single logical unit of work. Transactions are governed by ACID properties (Atomicity, Consistency, Isolation, Durability). Using START TRANSACTION, SAVEPOINT, and ROLLBACK, we can:

Perform operations safely,

Set a SAVEPOINT to mark a specific state,

ROLLBACK TO SAVEPOINT if a part of the transaction fails (e.g., duplicate primary key insertion), preventing data corruption.

#### 4. Error Simulation and Recovery

To test database robustness, simulated errors such as inserting duplicate enroll\_id values or invalid student\_id (violating FK constraints) help demonstrate:

How constraints enforce data validity,

How ROLLBACK TO SAVEPOINT can undo partial changes, leaving valid records intact.

#### 5. Data Retrieval with JOINS

Using JOIN operations, we can combine related data across multiple tables. For example:

A JOIN between Students, Courses, and Enrollments allows us to fetch each student's name, course title, and grade. This is essential for generating real-world reports and dashboards.

**6. Real-World Applications:-** This schema is similar to those used in:

College/University management systems,

Online learning platforms (e.g., Coursera, edX),

Training tracking systems in corporations.

- **CODES: -**

- Create Students table

- CREATE TABLE Students (  
    student\_id INT PRIMARY KEY,  
    name VARCHAR(100),  
    dob DATE  
);

- Create Courses table

- CREATE TABLE Courses (  
    course\_id INT PRIMARY KEY,  
    title VARCHAR(100)  
);

- Create Enrollments table with foreign key constraints

- CREATE TABLE Enrollments (  
    -- (Table structure continues with foreign key constraints)

```
enroll_id INT PRIMARY KEY,  
student_id INT,  
course_id INT,  
grade VARCHAR(2),  
FOREIGN KEY (student_id) REFERENCES Students(student_id),  
FOREIGN KEY (course_id) REFERENCES Courses(course_id)  
);
```

```
-- Insert into Students  
INSERT INTO Students VALUES  
(1, 'Ashish', '2002-03-14'),  
(2, 'Smaran', '2001-08-22'),  
(3, 'Vaibhav', '2003-01-05');
```

```
-- Insert into Courses  
INSERT INTO Courses VALUES  
(101, 'DBMS'),  
(102, 'Operating Systems'),  
(103, 'Computer Networks');
```

```
-- Begin Transaction  
START TRANSACTION;
```

```
-- First Enrollment: Ashish into DBMS  
INSERT INTO Enrollments VALUES (1, 1, 101, 'A');
```

```
-- Create SAVEPOINT  
SAVEPOINT sp1;
```

```
-- Second Enrollment: Ashish into Operating Systems  
INSERT INTO Enrollments VALUES (2, 1, 102, 'B+');
```

```
-- Third Enrollment: Ashish into Computer Networks  
INSERT INTO Enrollments VALUES (3, 1, 103, 'A');
```

```
-- Simulate an error: Duplicate enroll_id  
-- This will fail because enroll_id 2 already exists  
INSERT INTO Enrollments VALUES (2, 1, 101, 'C');
```

```
-- Rollback to savepoint to undo the faulty insert
```

```
ROLLBACK TO SAVEPOINT sp1;
```

```
-- Commit the transaction
```

```
COMMIT;
```

```
-- Final query to display student name, course title, and grade
```

```
SELECT
```

```
    s.name AS student_name,
```

```
    c.title AS course_title,
```

```
    e.grade
```

```
FROM Enrollments e
```

```
JOIN Students s ON e.student_id = s.student_id
```

```
JOIN Courses c ON e.course_id = c.course_id;
```

- **OUTPUTS: -**

	column_name name	data_type character varying	is_nullable character varying (3)
1	student_id	integer	NO
2	dob	date	YES
3	name	character varying	YES

Data Output Messages Notifications

	student_name character varying (100)	course_title character varying (100)	grade character varying (2)
1	Ashish	DBMS	A

- **LEARNING OUTCOMES: -**

- **Understanding of Database Schema Design:** Ability to create normalized relational tables using SQL.
- Understanding how to define **Primary Keys** and **Foreign Keys** to enforce entity integrity and referential integrity.
- **Application of Data Integrity Constraints:** Hands-on experience in applying **constraints** to prevent invalid or duplicate data entries (e.g., primary key violations, foreign key constraints).
- **Transaction Management in SQL:** Practical understanding of **ACID properties** in transactions.

- Use of START TRANSACTION, SAVEPOINT, ROLLBACK TO SAVEPOINT, and COMMIT for safe and controlled data operations.
- **Error Simulation and Recovery:** Ability to simulate errors (e.g., inserting duplicate keys) and recover using **ROLLBACK** techniques without losing valid data.